

NAME: K.Naga Venkata Siva

REGNO: 192425345

GEMINI:**1. Comparison of L1, L2, and L3 Cache Levels**

Parameter	L1 Cache	L2 Cache	L3 Cache
Typical Size	$32\text{ KB} - 64\text{ KB}$ per core	$512\text{ KB} - 1\text{ MB}$ per core	$8\text{ MB} - 64\text{ MB}$ (Shared)
Access Latency	$0.5 - 1.5\text{ ns}$ ($4 - 5\text{ cycles}$)	$3 - 7\text{ ns}$ ($12 - 14\text{ cycles}$)	$10 - 25\text{ ns}$ ($40 - 75\text{ cycles}$)
Bandwidth	Ultra-High ($>2 - 5\text{ TB/s}$)	High ($1 - 2\text{ TB/s}$)	Moderate-High ($300 - 800\text{ GB/s}$)
Throughput	2–3 memory ops / cycle	1 memory op / cycle	Multi-channel interleaved
Placement & Sharing	Inside core pipeline; private	Near core; private/semi-shared	Outside cores; shared across all cores

Impact on Performance	Eliminates immediate ALU stalls during instruction loops.	Reduces penalties from L1 data/instruction evictions.	Prevents high-penalty off-chip DRAM bus round-trips.
------------------------------	---	---	--

2. Differentiation of SRAM vs. DRAM

Parameter	SRAM (Static RAM)	DRAM (Dynamic RAM)
Cell Structure	6 Transistors (6T configuration)	1 Transistor + 1 Capacitor (1T-1C)
Access Latency	Ultra-low ($0.5 - 5 \text{ ns}$)	Moderate-High ($50 - 80 \text{ ns}$)
Bandwidth	Terabytes per second (on-chip buses)	$30 - 100 \text{ GB/s}$ per channel (off-chip bus)
Throughput	Immediate cycle-by-cycle pipelining	Limited by row-buffer activations/precharge cycles
Refresh Need	None (holds data while powered)	Periodic refresh required (causes latency jitter)
Density & Cost	Low density, physically large, high cost	Very high density, compact footprint, low cost
Hierarchical Role	L1, L2, and L3 CPU/GPU Cache memories	Main system memory (Primary working memory)

3. Virtual Memory with Paging vs. Cache Memory

Parameter	Cache Memory Mechanism	Virtual Memory with Paging
Management	100% Hardware-controlled (Cache Controller)	Software + Hardware (OS + MMU + TLB)
Granularity	Cache line (64 bytes)	Page (4 KB , 2 MB , or 1 GB)
Access Latency	$1 - 25\text{ ns}$	100 ns (TLB hit) to $5 - 10\text{ ms}$ (Page Fault from disk)
System Throughput	Multi-terabyte sustained throughput	Lower throughput; page faults stall the active thread
Address Mapping	Direct-Mapped / Set-Associative via tags	Multi-level Page Tables + TLB translation
Miss Penalty	Fetching line from DRAM ($50 - 100\text{ ns}$)	Disk I/O interrupt and context switch ($\mu\text{s} - \text{ms}$)

4. NAND Flash vs. DRAM

Parameter	NAND Flash	DRAM
Volatility	Non-Volatile (retains data without power)	Volatile (loses data on power loss)

Access Latency	High: $25 - 100\ \mu\text{s}$ (Reads), $200 - 500\ \mu\text{s}$ (Writes)	Low: $50 - 80\ \text{ns}$ (Reads and Writes)
Bandwidth	$2 - 7\ \text{GB/s}$ (over PCIe/NVMe Gen 4/5)	$50 - 100\ \text{GB/s}$ per DDR5 channel
Access Mode	Block/Page addressed (no byte access)	Byte-addressable
Endurance	Finite ($10^3 - 10^5$ write cycles per cell)	Virtually unlimited ($>10^{15}$ write cycles)
Hierarchical Role	Secondary Storage (SSDs, mass file storage)	Primary Memory (active execution space)

5. MRAM vs. RRAM (Emerging NVM Technologies)

Parameter	MRAM (Magnetoresistive RAM)	RRAM (Resistive RAM)
Storage Mechanism	Magnetic Tunnel Junction (MTJ resistance)	Conductive filament formation in dielectric
Access Time / Latency	Read: $1 - 10\ \text{ns}$, Write: $5 - 20\ \text{ns}$	Read: $10 - 50\ \text{ns}$, Write: $10 - 100\ \text{ns}$
Throughput	Very High (near-SRAM execution speeds)	High for sequential reads; moderate writes

Energy Efficiency	Extremely low active and zero standby power	Low write energy; ideal for Analog Matrix ops
Write Endurance	Extremely high (10^{12} - 10^{15} cycles)	Moderate to High (10^6 - 10^{10} cycles)
Primary Suitability	High-speed cache replacement, instant-on embedded	Dense AI model storage, Compute-In-Memory (CiM)

CHATGPT:

1. Comparison of L1, L2, and L3 Cache

Parameter	L1 Cache	L2 Cache	L3 Cache
Typical Size	32–64 KB per core	512 KB–1 MB per core	8–64 MB shared
Access Time	0.5–1.5 ns	3–7 ns	10–25 ns
Bandwidth	Very High	High	Moderate to High
Location	Inside CPU core	Close to CPU core	Shared among cores
Sharing	Private to each core	Private or semi-shared	Shared by multiple cores
Main Function	Stores immediately required data and instructions	Stores recently used data from L1	Stores larger shared working sets
Performance Effect	Minimizes CPU execution delays	Reduces L1 miss penalties	Reduces expensive DRAM accesses

Key idea:

L1 = Fastest and Smallest → L2 = Larger and Slightly Slower → L3 = Largest and Shared

2. SRAM vs. DRAM

Parameter	SRAM	DRAM
Cell Structure	Usually 6-transistor (6T)	1 transistor + 1 capacitor (1T-1C)
Access Time	Very low, typically below a few ns	Higher, typically tens of ns
Bandwidth	Extremely high for on-chip use	High for main-memory systems
Refresh	Not required	Periodic refresh required
Density	Low	High
Cost	Expensive	Less expensive
Power	Higher area and leakage per bit	Lower cost per bit but refresh consumes power
Typical Usage	CPU/GPU caches and buffers	Main system memory
Main Advantage	Very fast access	Large storage capacity

Key idea:

SRAM provides speed, while DRAM provides capacity.

3. Cache Memory vs. Virtual Memory Paging

Parameter	Cache Memory	Virtual Memory / Paging
Management	Mainly hardware controlled	OS + MMU + TLB
Data Unit	Cache line, commonly 64 bytes	Page, commonly 4 KB or larger
Purpose	Speeds up frequently accessed data	Provides large logical address space
Access Time	Usually a few to tens of ns	TLB hits are fast; page faults can take μ s to ms
Mapping	Tags and set-associative/direct-mapped structures	Page tables and TLB
Miss/Failure	Cache miss retrieves data from lower memory	Page fault may require storage access

Performance Impact	Small to moderate delay	Page faults can cause major delays
Storage Location	On-chip SRAM	DRAM with secondary storage backing

Key idea:
 Cache improves CPU speed, while virtual memory manages memory capacity and address translation.

4. NAND Flash vs. DRAM

Parameter	NAND Flash	DRAM
Type	Non-volatile	Volatile
Data Retention	Retains data without power	Loses data when power is removed
Access Speed	Relatively slow	Very fast compared with Flash
Typical Latency	Microsecond range	Tens of nanoseconds
Bandwidth	Lower than DRAM	Much higher
Access Method	Page/block based	Random byte/word access through memory interface
Endurance	Limited write endurance	Very high write endurance
Capacity	Very high	High
Typical Usage	SSDs and permanent storage	Main memory
Main Advantage	Large, persistent storage	Fast working memory

Key idea:
 NAND Flash stores data permanently, while DRAM stores data temporarily for fast processing.

5. MRAM vs. RRAM

Parameter	MRAM	RRAM
Storage Principle	Magnetic state of an MTJ	Resistance state of a memory cell
Type	Non-volatile	Non-volatile
Read Speed	Very fast	Fast
Write Speed	Fast	Moderate to fast, depending on technology
Power Consumption	Low standby power	Low power, especially for suitable workloads
Endurance	Very high	Generally lower than MRAM, depending on cell technology
Density	High	Potentially very high
Main Strength	Fast, durable, persistent memory	Dense storage and in-memory computing potential
Typical Application	Embedded memory, cache replacement, instant-on systems	AI accelerators, model storage, Compute-in-Memory
Main Advantage	High speed and endurance	High density and efficient computation near/in memory

Key idea:

MRAM focuses on speed and endurance, while RRAM is especially promising for dense storage and AI Compute-in-Memory applications.